

# Mining Jazz: An Experience Report

Thanh H.D. Nguyen, Adrian Schröter, Daniela Damian  
Software Engineering Global interAction Lab (SEGAL)  
Department of Computer Science  
University of Victoria  
Victoria, British Columbia, Canada  
{duythanh, schadr, danielad}@uvic.ca

## ABSTRACT

Integrated collaborative software engineering tools such as IBM Jazz<sup>TM</sup> provide researchers opportunities to access a vast amount of information that reflects the real development activities of software teams. With all of the artifacts stored in one place and linked together, researchers can gain valuable insights about the teams' collaborative activities. We report on our experience in mining the Jazz development team repository to study collaboration in software development. In particular we explain the different options and constraints we had during our efforts to mine the Jazz repository. In addition, we show some results we obtained by studying communication practices in the Jazz project team to demonstrate our research with data obtained from the mining process.

## 1. INTRODUCTION

Accessing research data in Software Engineering is very difficult. Regardless of the project methodology, the high cost of labour and developers' stressful workload are always big challenges that researchers face in their effort to access developers' working environment. On the other hand, software developers usually utilize digital media to coordinate their work. This media is usually archived and thus can be used for research as it provides valuable insights into the developers' activities. For example, the bugzilla bug tracking database of several open-source development teams have been studied in the past [3]. So did developer mailing lists, emails [4], or even IRC chat [5].

However, this media are usually not connected to each other. It is very difficult for researchers to link artifacts such as a bug on bugzilla and a post on the developer mailing list to make them useful for research. Fortunately, the advent of integrated end-to-end collaboration development environments such as IBM Jazz [1] provides research with opportunities to explore the development activities without having to take away developers' valuable working hours or manually connecting different artifacts from different media. In this paper, we are describing our experience in data mining the IBM Jazz team collaboration. In the next section, we present our conceptualization of doing research via artifacts as well as challenges we encountered. Then we explain the importance of the IBM Jazz development environment to research. In Section 4, we explain the different options and constraints we have during the mining process. Section 5 outlines the three design principles we learned from building the data mining tool. And finally, in Section 6, we provide some

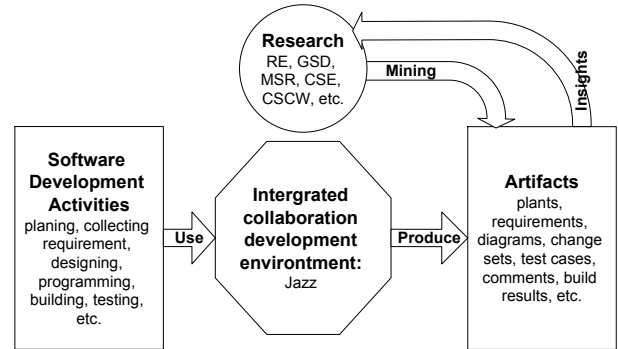


Figure 1: The research via artifacts process

results of a research project we did based on the data we mined in Jazz.

## 2. RESEARCH VIA ARTIFACTS

Our research methodology applied and extended the growing knowledge of mining software repositories [9]. Figure 1 illustrates the process of mining and leveraging software artifacts for research as we conceptualize it. The goal is to study different kinds of *Software Development Activities*, such as planning, designing, implementing, or testing a product. With the use of *integrated development environments*, the activities are recorded. In highly integrated collaboration development environments such as Jazz, the record of development activities are kept in the form of different artifacts such as plans, work items, or source code changes. We opted to infer knowledge from the created artifacts, because gathering data directly from the developers, through surveys or interviews, are often highly invasive and very time consuming for researchers.

The mining and studying of development artifacts has its own specific challenges as well. Two main challenges are:

1. The complexity of the data mining task.
2. The validity of insights obtained through the artifacts.

The first challenge is a technical one. It raises questions such as "How to get access to the data?". The data mining process can be very complex and sometimes impossible. As we

will explain in subsequent sections, it requires careful planning and implementation. The second challenge is on the research level. By doing research via artifacts, we have to make inferences about the software development activities through the artifacts it produces. This makes it hard to interpret the result. The main questions center around "How to interpret the artifacts?" and "Is our interpretation valid?".

The importance of these challenges depends on the point of view. For the development team the first challenge is more relevant because extracting data from a repository might crash it or delay the development process in other ways. Whereas the second challenge is of more relevance to research. Researchers are always interested in new findings, growing, or validating the body of knowledge. After leveraging the two point of views, the big question that we are trying to answer is: How to build a Jazz data mining tool that minimizes the effort of both the development team and the researchers, and maximizes the validity of our study? Before we discuss our approach to addressing this question, we briefly describe the IBM Jazz project in the next section and explain why we decided to build the data mining tool.

### 3. WHY JAZZ

Jazz [6, 7] is a development environment that integrates programming, communication, and project management. The main goal of Jazz is to bring collaboration support for all of the software development activities such as planning, building, testing, and reporting. For developer, this means that they can communicate and coordinate their work through the very same tools they use to develop code. To achieve this, Jazz provides a central repository for all tools supporting the development process such as a planning wiki, a work item tracking system, a report component, and a build system.

At the time of our study, Jazz consisted of three major components [1]:

- The Jazz repository.
- The Rational Team Concert (RTC) client.
- The dashboard, a web user interface.

The *Jazz repository* stores all artifacts available to the two clients. Examples for such artifacts are work items, comments, and build results. The *RTC client* is the main development client. It is built upon the Eclipse platform and allows developers to control their work spaces, to perform version control functions, and to request builds. The *web user interface* allows developers to access most of the non-coding related artifacts, such as work items, iteration plans, and reports.

For researchers, Jazz provides an easy way to extract valuable research data. Due to the tight integration of the different components, links of artifacts to related software development activities already exists. For example, a comment is linked to tasks, also called work items, and source code changes in the form of change sets are attached to these work items. This is a major advancement to other systems that used their own heuristics to link software artifacts [11].

## 4. OPTIONS AND CONSTRAINTS

The Jazz repository gave us several choices on how to extract data from it: via the dataware house report, via the RTC client API, via the server API, and via the back-end database. Each of these choices comes with its own constraints. These constraints can be summarized in three categories: (1) the *effort* needed, (2) the *logistical overhead* in terms of travel or non legal permissions, and (3) the *legal* issues such as guaranteeing the privacy of developers. In addition to these constraints the four choices provided by the Jazz repository differ in the system level they are deployed at.

Option	Effort	Logistic	Legal
Report	Hard	Easy	Easy
Client API	Medium-Hard	Medium	Easy-Medium
Server API	Medium-Hard	Medium-Hard	Medium
Database	Easy	Medium	Hard

**Table 1: Our simple risk analysis based the available options and the constraints**

**Report** The report component in Jazz caches performance statistics of the Jazz repository and provides mechanisms to assess this information via the Jazz Dashboard. We could build a report that queries the data from the Jazz team server.

The draw back in terms of effort following this option is that we need to learn how to set up the report and obtain the data from the given data warehouse. Note that the report component had changed since then. In term of effort, this option became easier now. Logistically, this is the easiest option because the report can be deployed and obtained online. It is also the easiest option in term of legal requirement because all of the information in the dataware house are supposed to be public.

**Client API** Using the RTC client API we can write a plug-in to access the data via the RTC client. This plug-in can query the data that is available to the client and serialize the result. We can then import the result into a relational database.

The effort we need to spend in developing a plug-in using the client API lies in learning it. Although we are familiar with Java and Eclipse technologies, the APIs were constantly changed. More importantly, we are also restricted to the queries that are available to the client. We have to search the client code to find suitable queries and work around their restrictions. In terms of logistic, the development can be done off-site. The deployment have to be done on-site but it is not difficult because the query can be run inside a standard developer set up. Using the client API does not require a lot of legal advise because most data that is available through the client is available to the public.

**Server API** Another option to query the data is bypassing all the client APIs and access the Jazz repository itself. This way we are still restricted by the repository API but not by the client API. Note that these APIs are not formally defined yet at the time of writing this article.

Since we can bypass the client API, this option would take less effort on the programming. But it requires deep understanding of what the repository can provide. The learning process would have take as much effort as the client API option. In term of logistic, we would be able to run it on-site and it is harder to debug off-site because it can not be run from a nornal client. It may require working on-site with the repository team. Legally, this option can be tricky because we could get information that are not normally available to the public.

**Database Backend** If we are allowed, we can also acquire a copy of the backend database of the Jazz repository.

This option would be the least effort in term of programming. However, it would be very difficult, in term of legal requirement, to warranty privacy and intellectual property of the compant. Even if we are allow to, we would require to manage the task of getting the large data set and extracting the useful information from the raw data.


Table 1 shows the simple risk analysis we did based on the options available and the constraints. We decided to obtain the data via the RTC client API. This option seem to have the average of medium risk. In term of effort, it is very time consuming but we have the resources to manage. For us, the logistic and legal requirement are more important and they are manageable on this option.

## 5. DESIGN PRINCIPLES

The main requirement for our tool besides being able to extract all data is to not cause any problem in the development cycle of the team, such as causing to much load on their repository while extracting the data. We implemented these requirements in a program that was developed in several iterations.

*Iterative.* At the beginning of developing our mining tool, we evaluates the data that was available through the RTC client API. We found that the API did not support the extraction of all the data so we proceeded with an iterative development that was oriented at the different kind of data accessible through the API at a time. Similar to iterations in an Agile process, we extended the extraction capability of the tool one by one via adding data query set. At the end of each iteration, we asked our research contact at IBM to extract more data with the updated tool. This had the advantage that we received additional feedback about API usage from our IBM research contact.

*Minimally Invasive.* The data quantity we want to extract can cause a high load on the repository and thus disrupt the development process, each time we extract the data. We avoided causing such high repository load by updating the existing data instead of extracting all data every time we run the program. This had the advantage in case we finished a new query we did not need to extract all the data we extracted earlier again.



**Figure 2: The communication based social network of the Jazz development team**

*Separation.* In order to keep the additional load minimal on the repository we tried to remove redundancies in the queries as well as avoid the queries to load data that has already been extracted. At the beginning, some of our queries depended on each other or we could not avoid loading already extracted data. For example, if we wanted to fetch the new comments on work items, we had to fetch the work items themselves although they had not changed since the last extraction. This introduced additional load on the repository. We removed all unnecessary dependencies among queries and tried to minimize loading data that has already been extracted.

## 6. COMMUNICATION DELAY PROJECT

This project [10] used data we mined as described in this paper and we describe it as an illustration of our design principles and conducting research using software development artifacts.

There were two main objectives of this project: (1) to examine if the Jazz team experience communication and task completion delay due to the large geographic distance between the teams and (2) gain insights about communication patterns of the entire project team. For the first objective, we conceptualized communication delay as the difference between comments on work items and task completion delay by the resolution time of a work item. Thus, we had to query all work items from the Jazz repository. With the information on work items we addressed our second objective by constructing and examining a project-wide communication-based social network that connecting developers that commented on the same work items.

Following an iterative process, we first designed a query for a teams' hierarchy. Subsequently, we wrote a query to retrieve the members of each team. In the next iteration we designed the query to retrieve all work items. This query induced a high load on the repository due to the large amount

of data that was contained in all work items. During the early versions of Jazz we encountered client and server side failures cause by our query. To minimize the additional load generated from re-running a crashed query, our tool stores a cache of the queried data which does not need to be extracted during the re-run. When it is resumed, the query picks up from where it left. We were able to extend the tool and extract the data within the first six months of our project.

With the information from the work items, our analysis suggests that the Jazz team did not experience as much delay as reported in previous literatures [2, 8]. We hypothesize that this is due to the new integrated development support provided by Jazz as well as advances in software practices used by the Jazz distributed team.

Using this data we also constructed the project-wide social network based on communication as shown in Figure 2. The shape of the nodes (and the flower-like petals) indicate the developers' geographical location. The color of the nodes indicate how active each developers is in the overall network from most active (red) to least active (light blue). The red circle emphasizes the active core members of the project.

## 7. CONCLUSION

Mining the Jazz repository has been a challenging journey. In this paper we shared our experiences and the lessons we learned during the process: iterative, minimal invasive, and separation. We think that the design principles we presented here are applicable to any research via artifact projects regardless of the repository behind it.

For us, the Jazz data continue to be a very valuable source of information. The artifacts in Jazz hold valuable insights about both the social and technical aspects of the Jazz team's development process. Obtaining the communication and coordination data in one place creates numerous of possibilities for researchers to look at various aspects of intricate software development activities all together, and which were very hard to do with research methods other than repository mining. With the release of Jazz and the adoption of technology by many of the Eclipse based products from IBM, we think that mining the Jazz repositories will bring more and more benefits to research.

## 8. ACKNOWLEDGMENTS

We want to thank all members of the IBM Jazz project who helped us, particularly Harold Ossher, Li-Te Cheng of IBM TJ Watson, and Jean-Michel Lemieux of IBM Ottawa Lab. Many thanks to the members of the SEGAL lab. Funding was provided by NSERC and IBM Jazz Technology Fellowship Award.

IBM and Jazz are trademarks of IBM Corporation in the US, other countries, or both.

## 9. REFERENCES

- [1] IBM Rational Jazz. <http://jazz.net/>, 09 2008.
- [2] R. D. Battin, R. Crocker, J. Kreidler, and K. Subramanian. Leveraging Resources in Global Software Development. *IEEE Software*, pages 70–77, March/April 2001.
- [3] N. Bettenburg, S. Just, A. Schröter, C. Weiss, R. Premraj, and T. Zimmermann. What makes a good bug report? Technical report, Universität des Saarlandes, Saarbrücken, Germany, September 2007.
- [4] C. Bird, A. Gourley, P. Devanbu, M. Gertz, and A. Swaminathan. Mining email social networks. In *MSR '06: Proceedings of the 2006 international workshop on Mining software repositories*, pages 137–143, New York, NY, USA, 2006. ACM.
- [5] M. Cataldo, P. A. Wagstrom, J. D. Herbsleb, and K. M. Carley. Identification of coordination requirements: implications for the Design of collaboration and awareness tools. In *CSCW '06: Proceedings of the 2006 20th anniversary conference on Computer supported cooperative work*, 2006.
- [6] L.-T. Cheng, C. R. B. de Souza, S. Hupfer, J. Patterson, and S. Ross. Building Collaboration into IDEs. *ACM Queue*, 1(9), 12 2003.
- [7] R. Frost. Jazz and the Eclipse Way of Collaboration. *IEEE Software*, 24(06):114–117, 2007.
- [8] J. D. Herbsleb and A. Mockus. An Empirical Study of Speed and Communication in Globally Distributed Software Development. *IEEE Transactions on Software Engineering*, 29(6):481–494, June 2003.
- [9] M. Lanza, M. W. Godfrey, and S. Kim. Msr 2008 - 5th working conference on mining software repositories. In *ICSE Companion '08: Companion of the 30th international conference on Software engineering*, pages 1037–1038, New York, NY, USA, 2008. ACM.
- [10] T. Nguyen, T. Wolf, and D. Damian. Global software development and delay: Does distance still matter? In *Proceedings of the 3rd International Conference on Global Software Engineering (ICGSE 2008)*. IEEE Computer Society, August 2008.
- [11] D. Čubranić and G. C. Murphy. Hipikat: recommending pertinent software development artifacts. In *ICSE '03: Proceedings of the 25th International Conference on Software Engineering*, pages 408–418, Washington, DC, USA, 2003. IEEE Computer Society.